



# Using COMTiles to reduce the hosting costs of large map tilesets in the cloud

---

FOSS4G 2022

Markus Tremmel



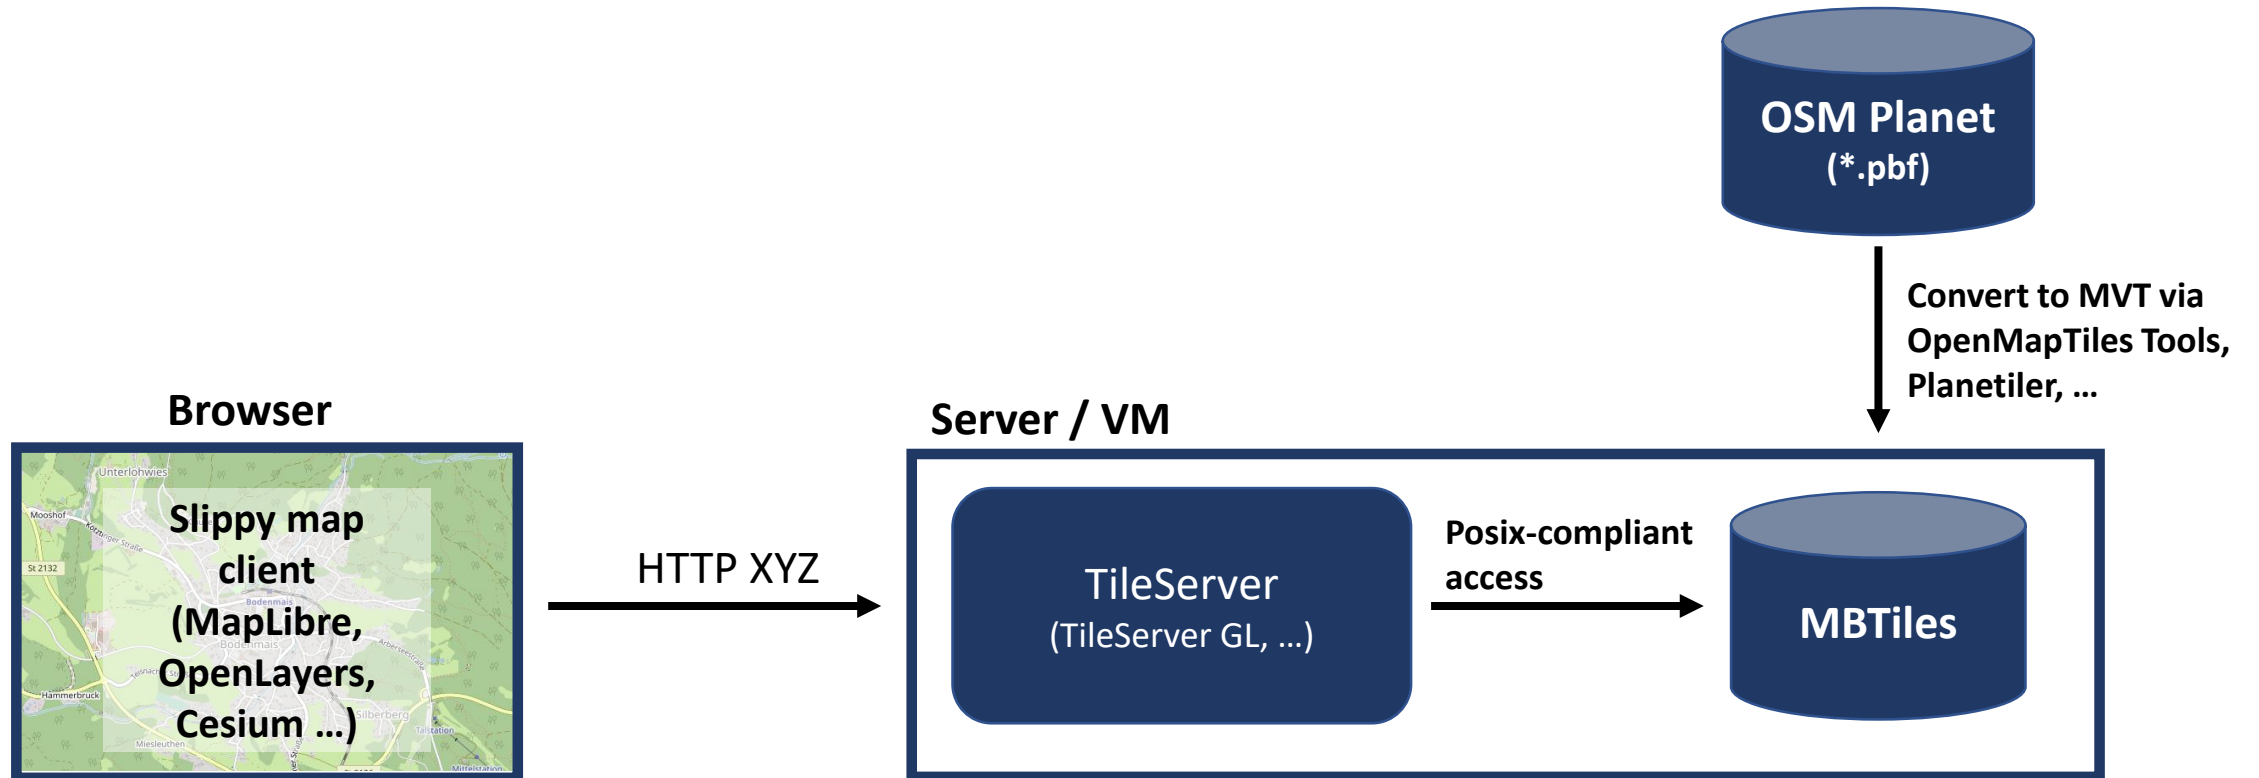


# About me

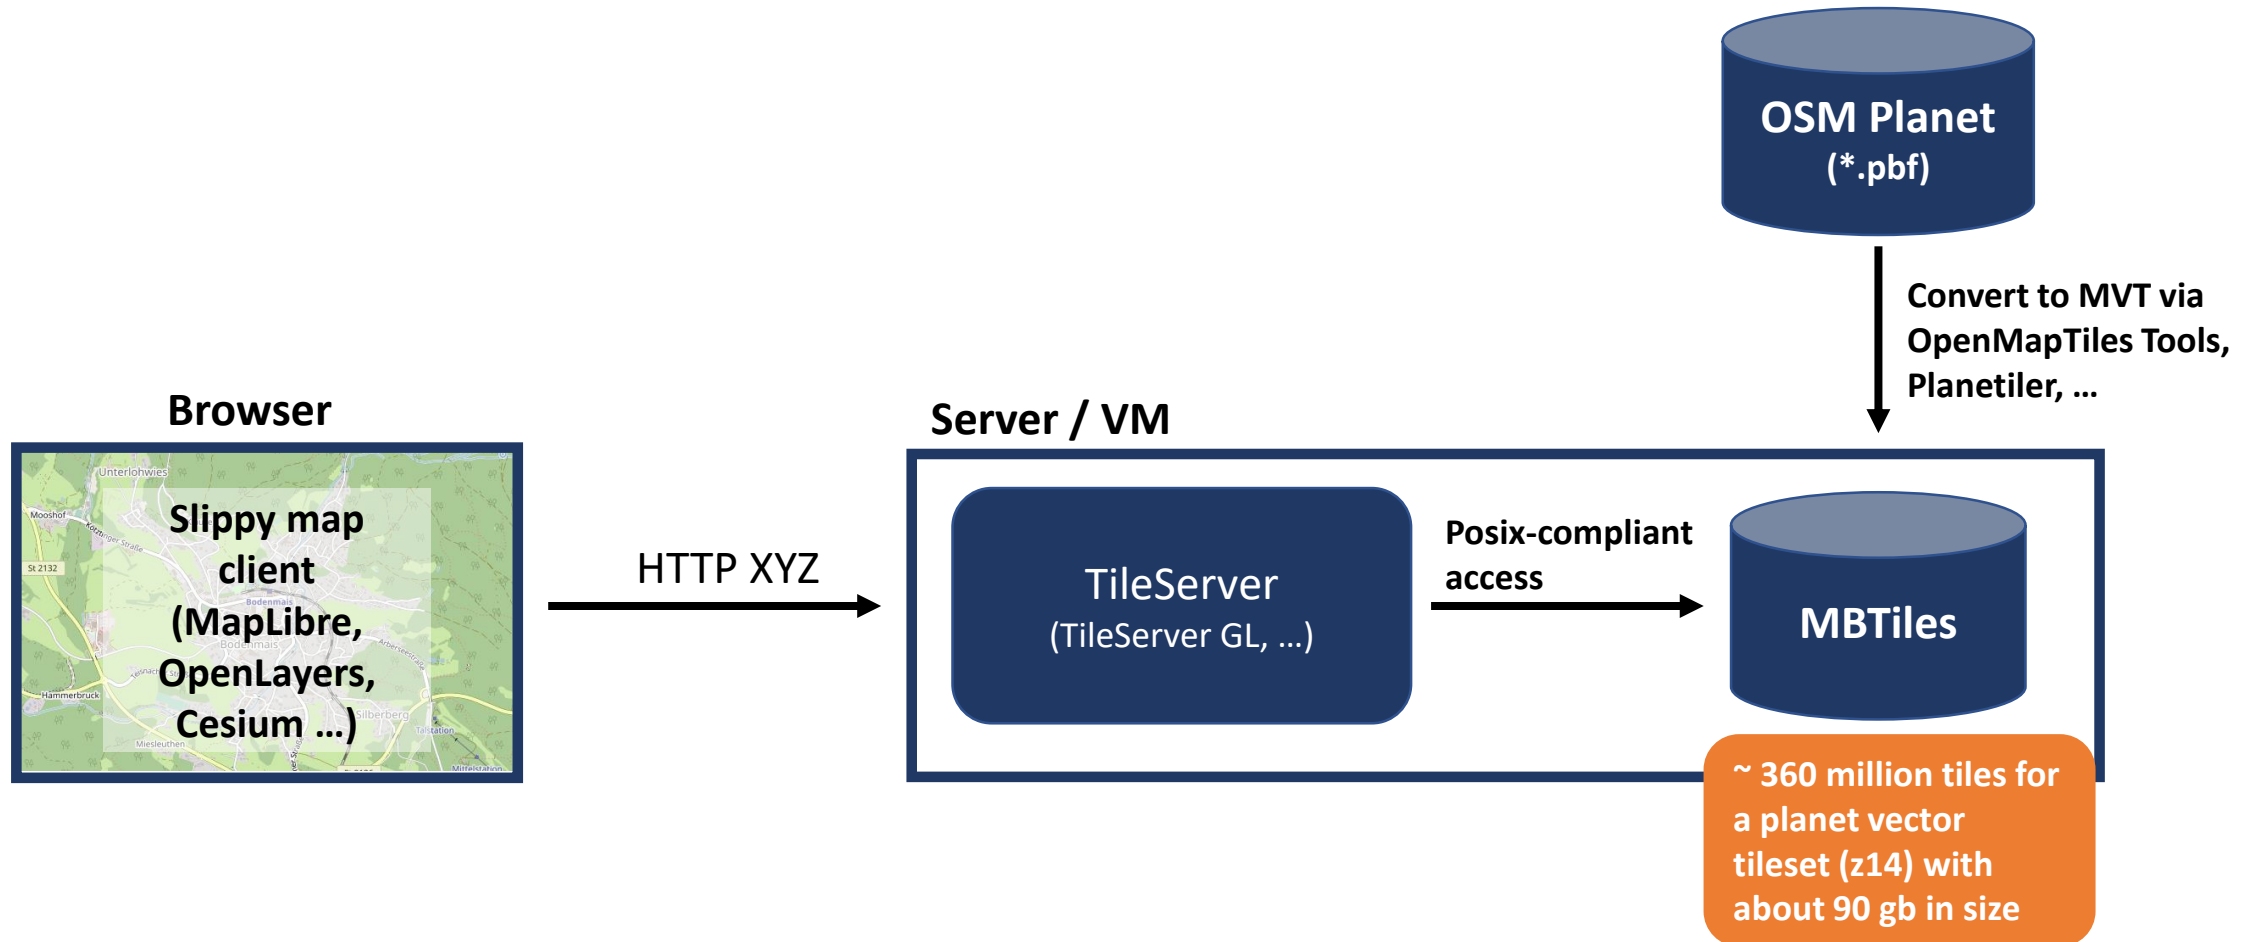
- Geospatial Software Architect at Rohde and Schwarz in Germany
- Lecturer at the Deggendorf Institute of Technology (DIT) in applied computer science with focus on spatial systems
- Main (research) interests: Cloud Native Geospatial, 3D Maps (3DTiles), Map Renderer



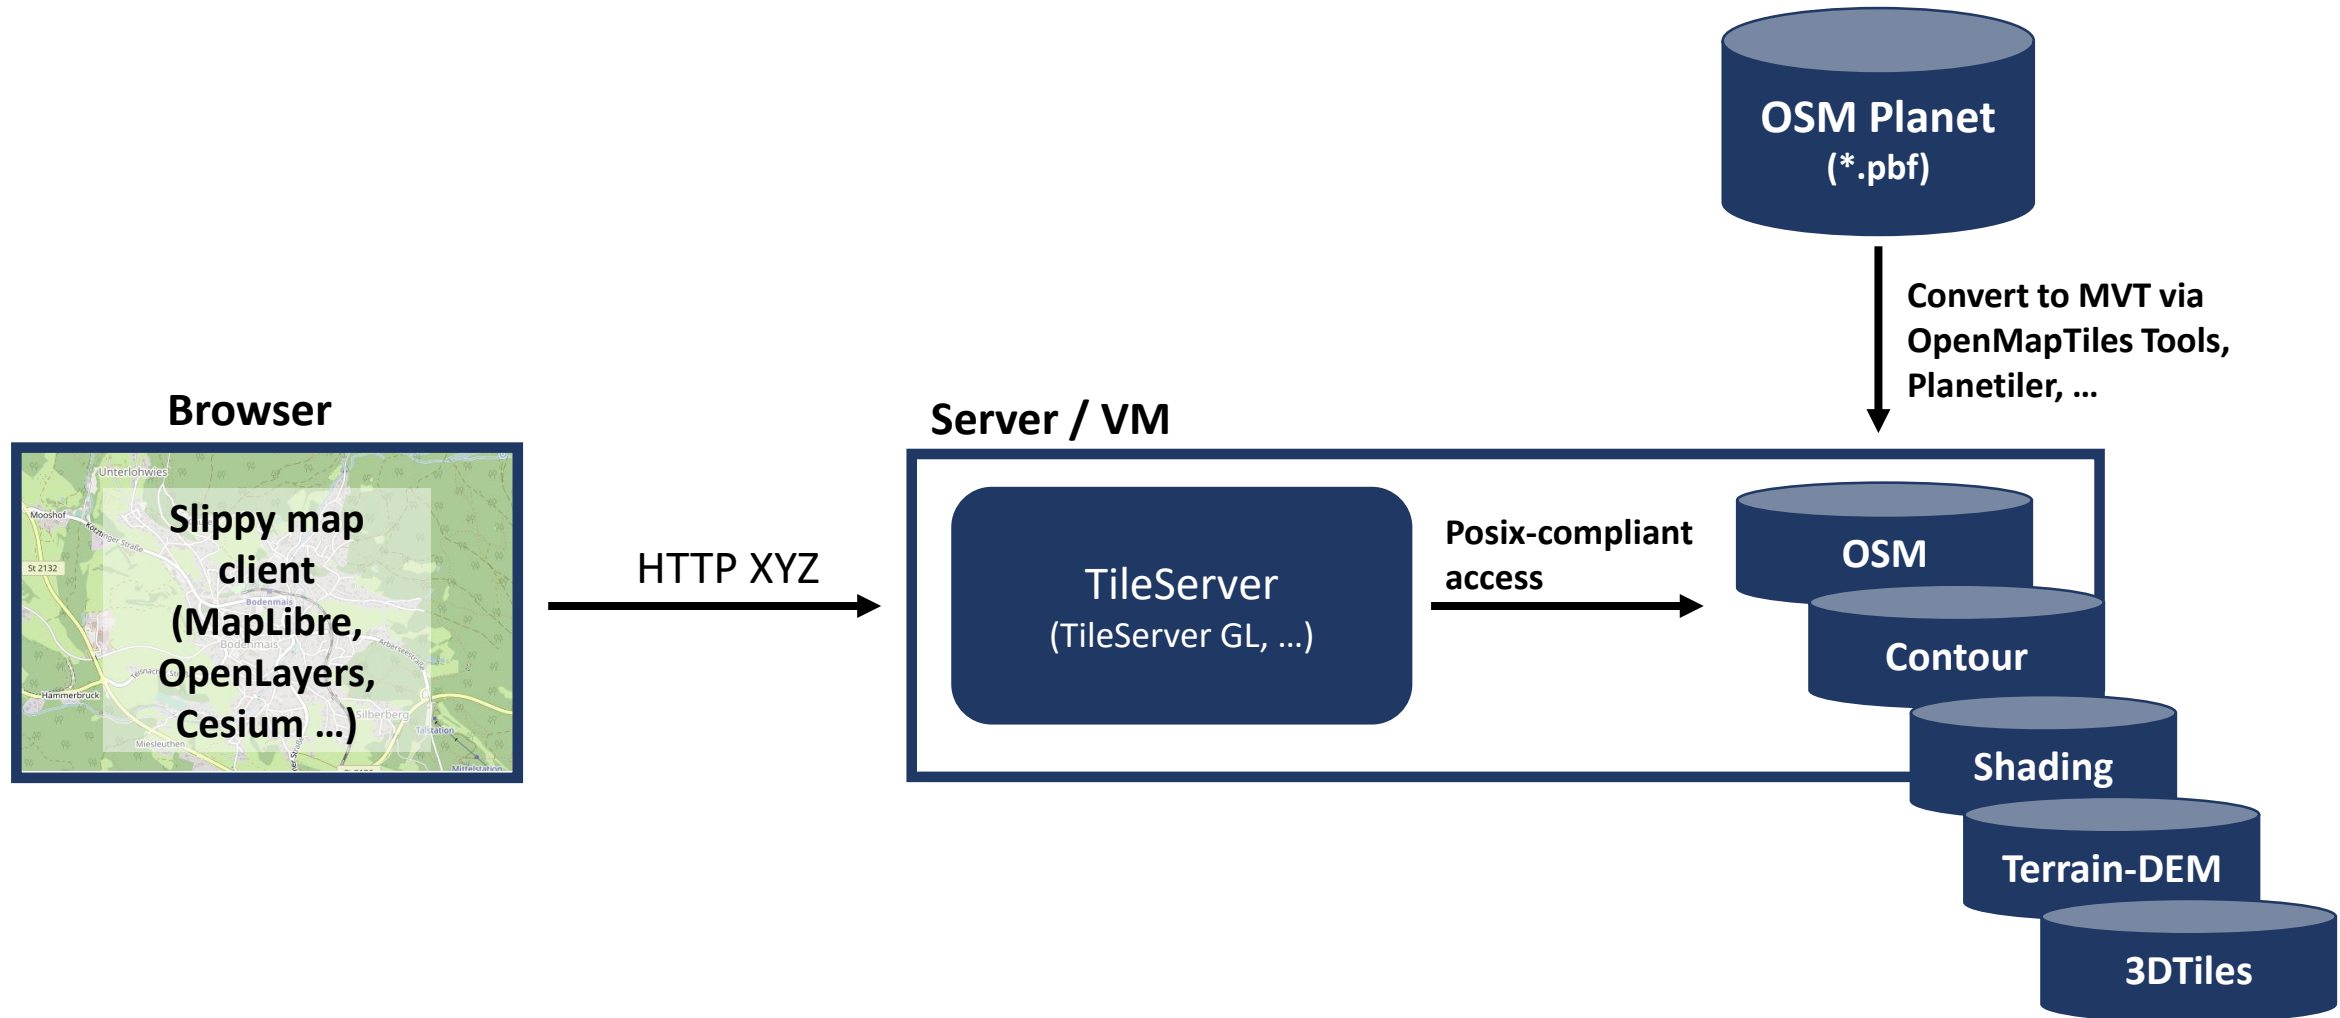
# Conventional deployment of map tilesets



# Conventional deployment of map tilesets



# Conventional deployment of map tilesets

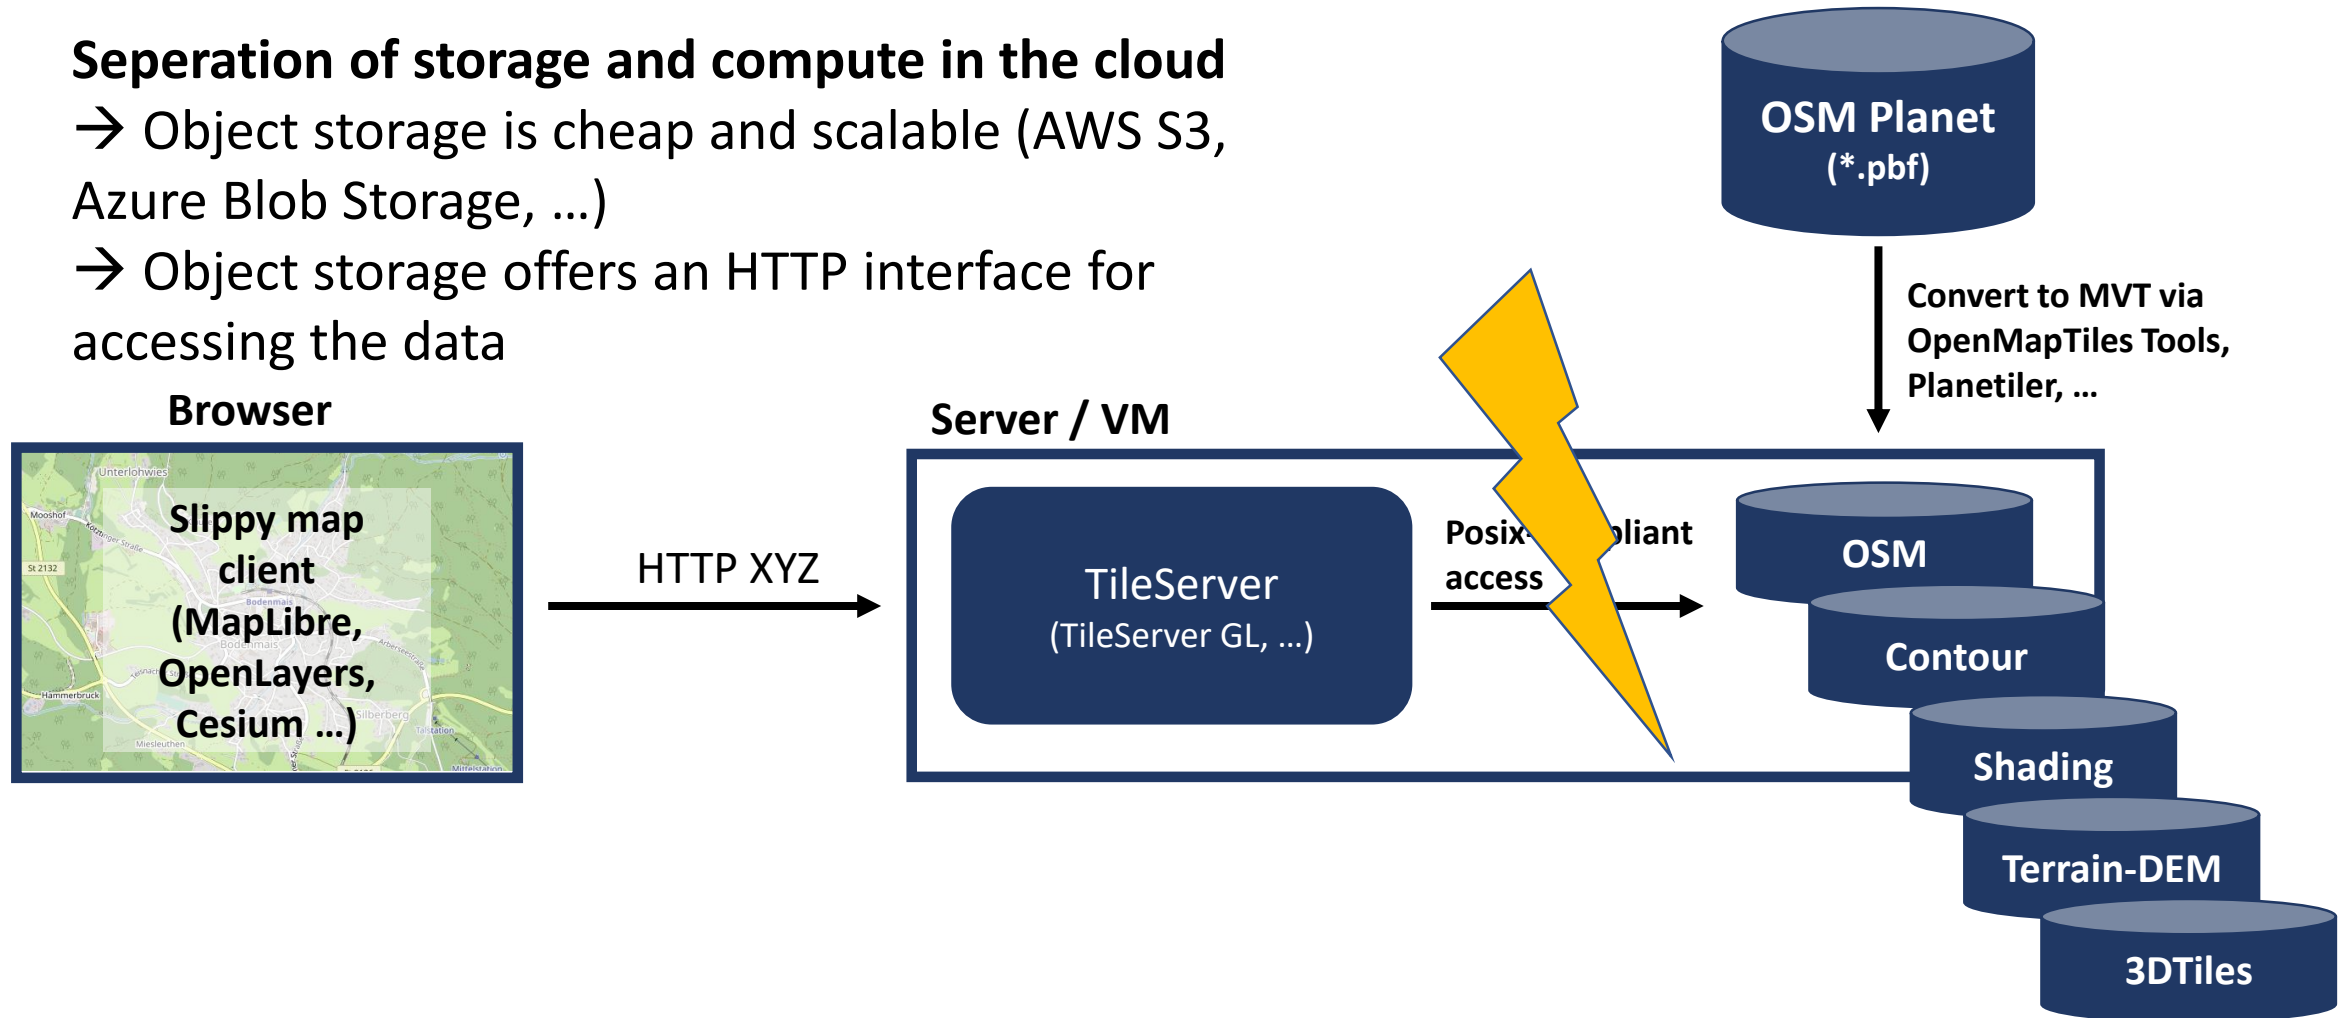


# Conventional deployment of map tilesets

## Seperation of storage and compute in the cloud

→ Object storage is cheap and scalable (AWS S3, Azure Blob Storage, ...)

→ Object storage offers an HTTP interface for accessing the data

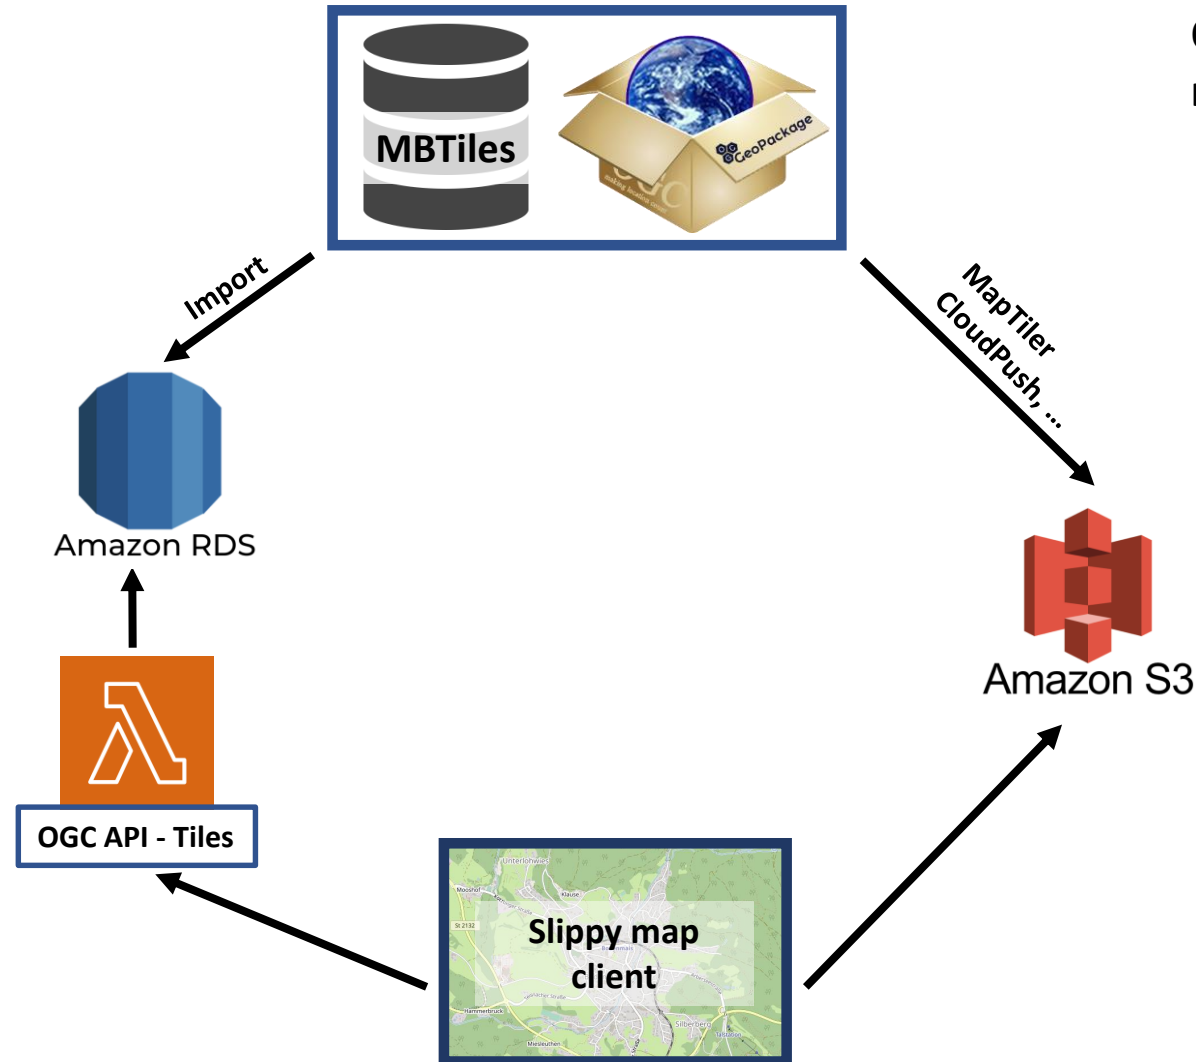




# Options for deploying map tilesets in the cloud

**Option 1: backend with database and tileserver**

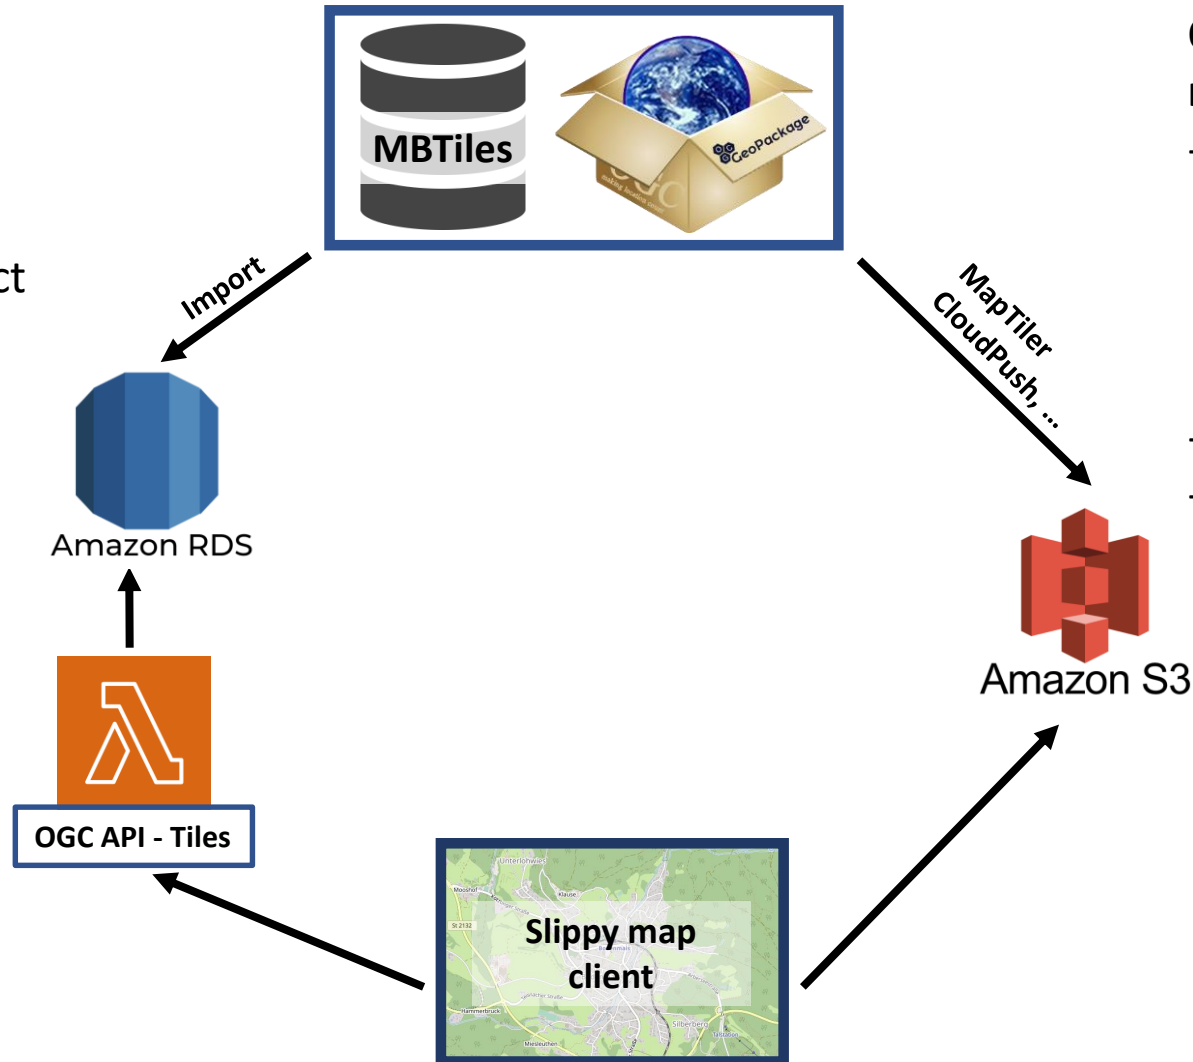
**Option 2: directly hosting the map tiles on an object storage**



# Options for deploying map tilesets in the cloud

## Option 1: backend with database and tileserver

- significantly more expensive compared to hosting only on an object storage
- ...



## Option 2: directly hosting the map tiles on an object storage

- Loading up hundreds of millions of tiles to an cloud storage for a planet scale vector tiles dataset is expensive
- Many repeated (raster) tiles
- ....



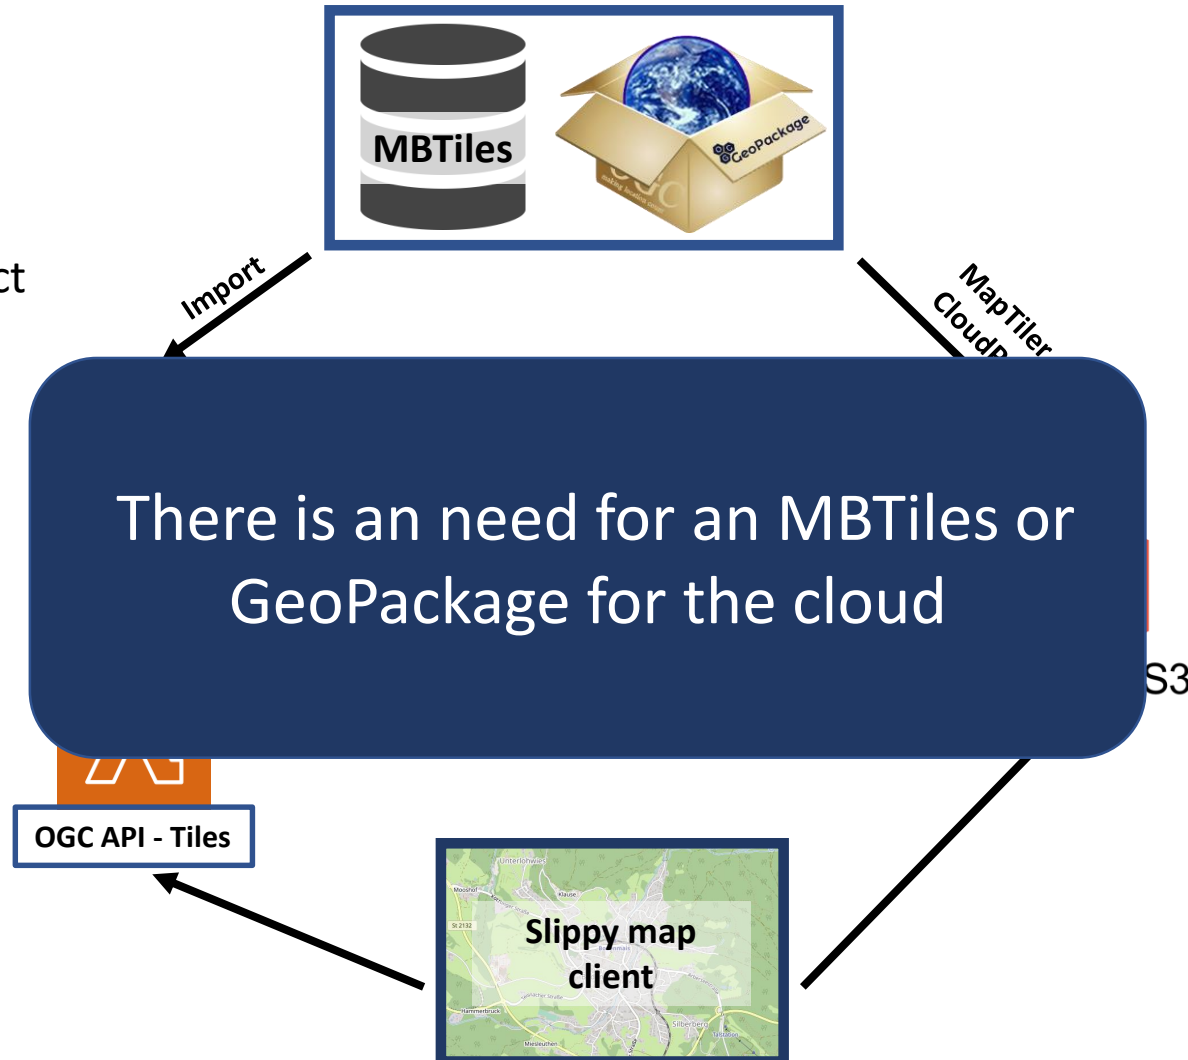
# Options for deploying map tilesets in the cloud

## Option 1: backend with database and tileserver

- significantly more expensive compared to hosting only on an object storage
- ...

## Option 2: directly hosting the map tiles on an object storage

- Loading up hundreds of millions of tiles to an cloud storage for a planet scale vector tiles dataset is expensive
- Many repeated (raster) tiles
- ....



# Categories of Cloud Optimized Formats

## **Raster (Satellite)**

Cloud Optimized GeoTiff (COG)

## **Point Clouds**

Cloud Optimized Point Cloud

## **Vector**

FlatGeobuf  
GeoParquet

## **N-dimensional Arrays**

Zarr  
TileDB

## **Tile Archives**

PMTiles  
COMTiles

# Categories of Cloud Optimized Formats

## **Raster (Satellite)**

Cloud Optimized GeoTiff (COG)

## **Point Clouds**

Cloud Optimized Point Cloud

## **Vector**

FlatGeobuf  
GeoParquet

**Focus on Analysis**  
-> (heavy) analytical  
workflows/processing

## **N-dimensional Arrays**

Zarr  
TileDB

## **Tile Archives**

PMTiles  
COMTiles

**Focus on  
Visualization**  
-> “Slippy map user  
experience”

# Cloud native principles



Objects storage like AWS S3




Read portions (chunks) of a file  
via HTTP range requests



(streamable) spatial index



Metadata for describing the  
tileset



**Request: planet.comt**  
Range: bytes = 0 – 1000

**Response**  
206 Partial Content  
Content-Range: bytes 0-1000

planet.comt = 90gb



Amazon S3

# COMTiles

---

Based on the ideas of Cloud Optimized GeoTIFF and extended for the usage of raster and in particular vector map tilesets

---

COMTiles are a streamable and read optimized file archive for hosting map tiles at global scale on a cloud object storage

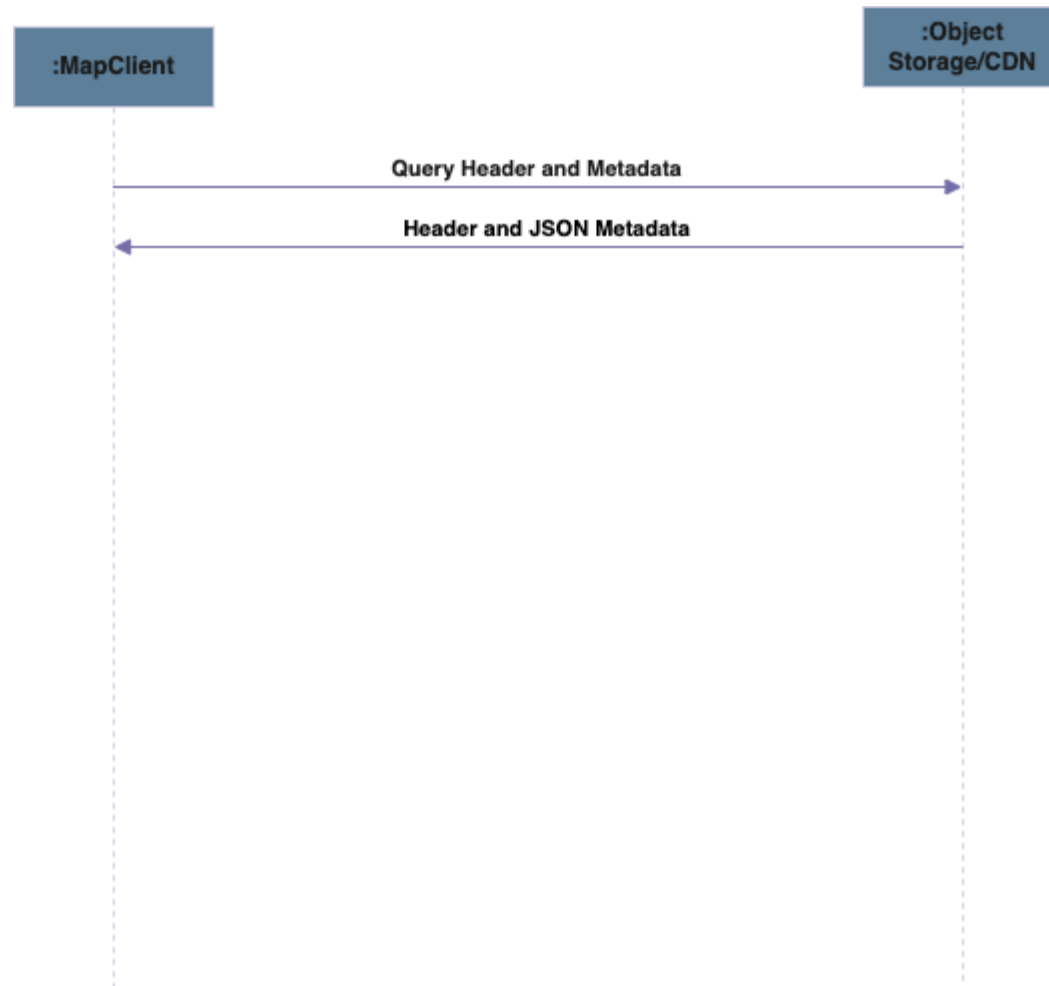
---

The map tiles can be accessed directly from a browser via HTTP GET range requests

---

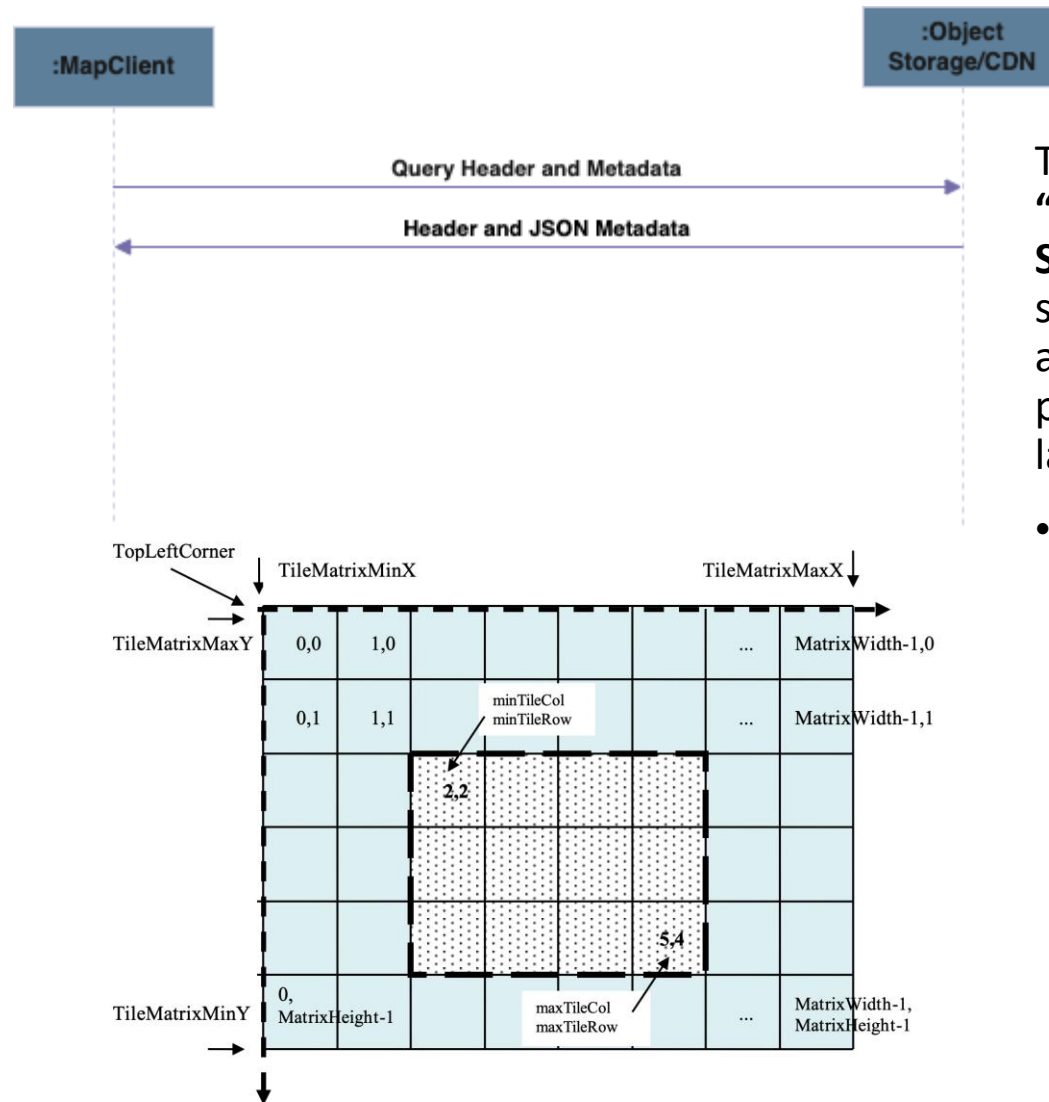
The main focus of COMTiles is to significantly reduce costs and simplify the hosting of large raster and vector tilesets at global scale in the cloud because no backend (database, server) is needed

# COMTiles sequence





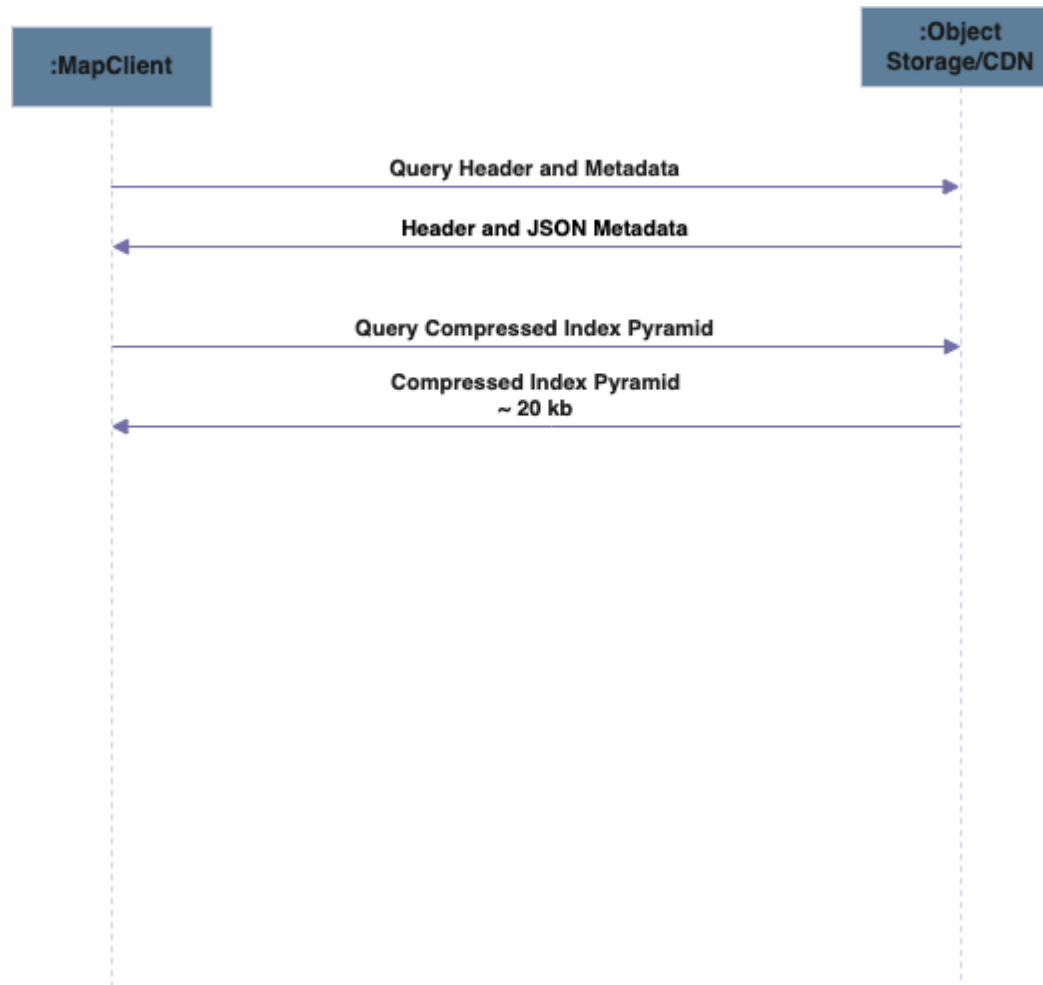
# COMTiles sequence



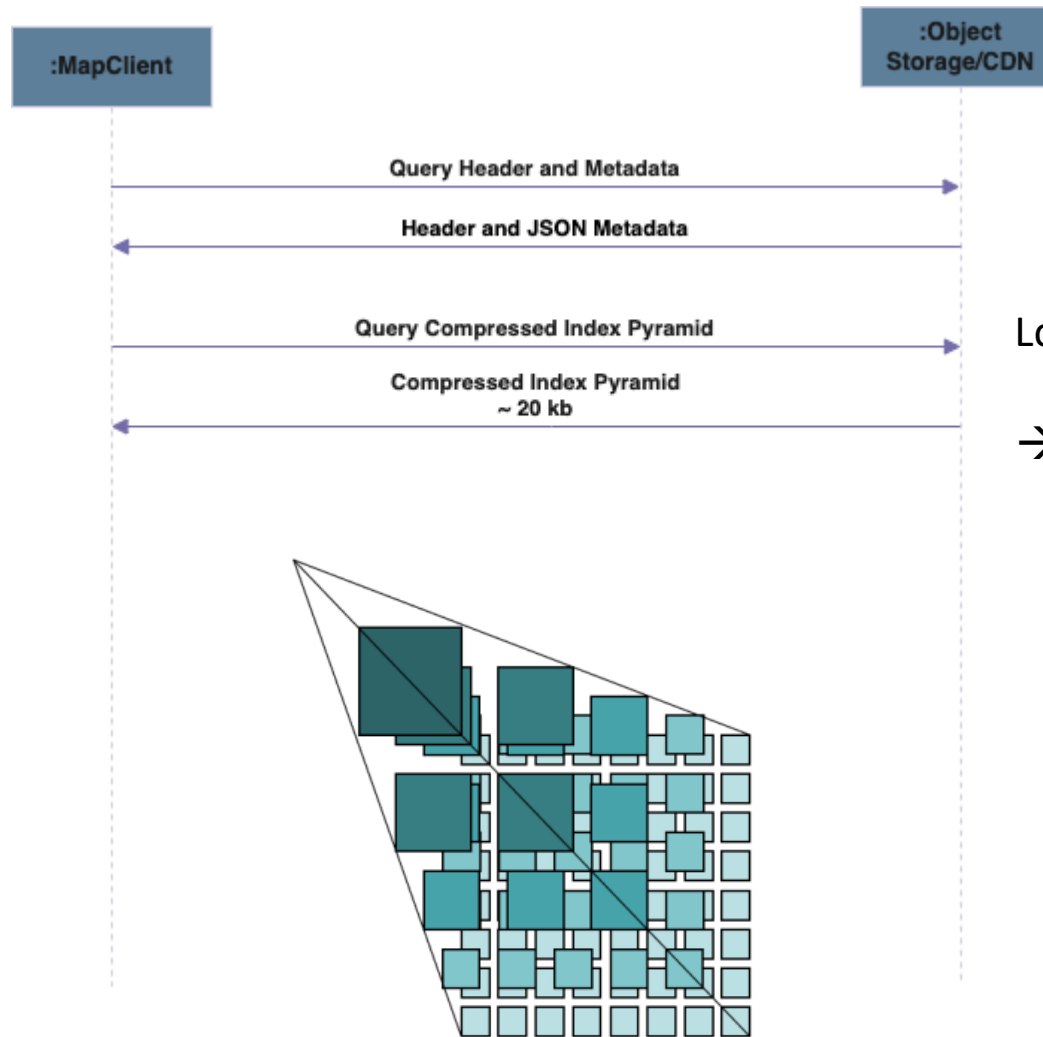
The metadata are based on the **“Two Dimensional Tile Matrix Set and Tile Set Metadata”** specification (currently draft) and extend it with additional properties about the index layout

- Concepts
  - **TileSet metadata** -> attribution, crs, layers, dataType, ...
  - **TileMatrixSet** -> Common TileMatrixSetDefinitions like WebMercatorQuad
  - **TileMatrixSetLimits** -> minTileRow, maxTileRow, minTileCol, maxTileCol, ...

# COMTiles sequence



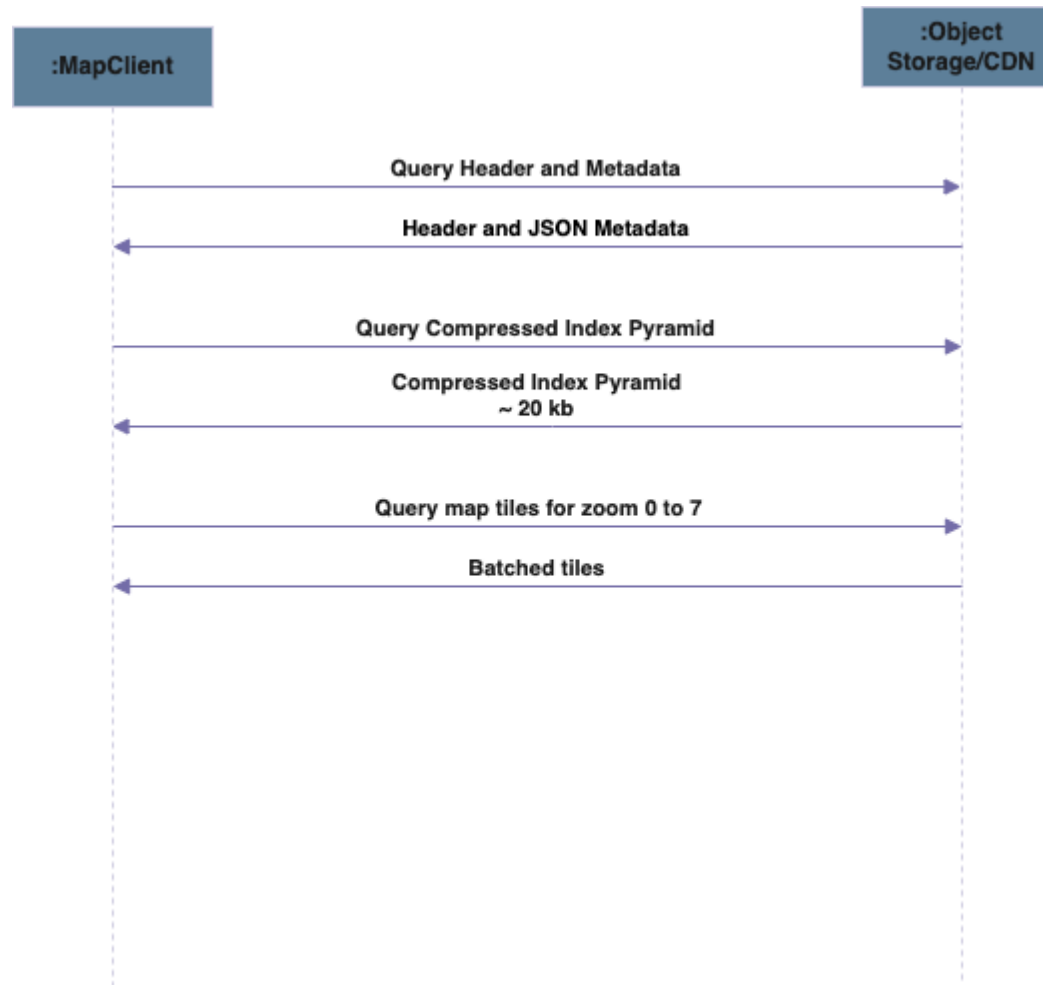
# COMTiles sequence



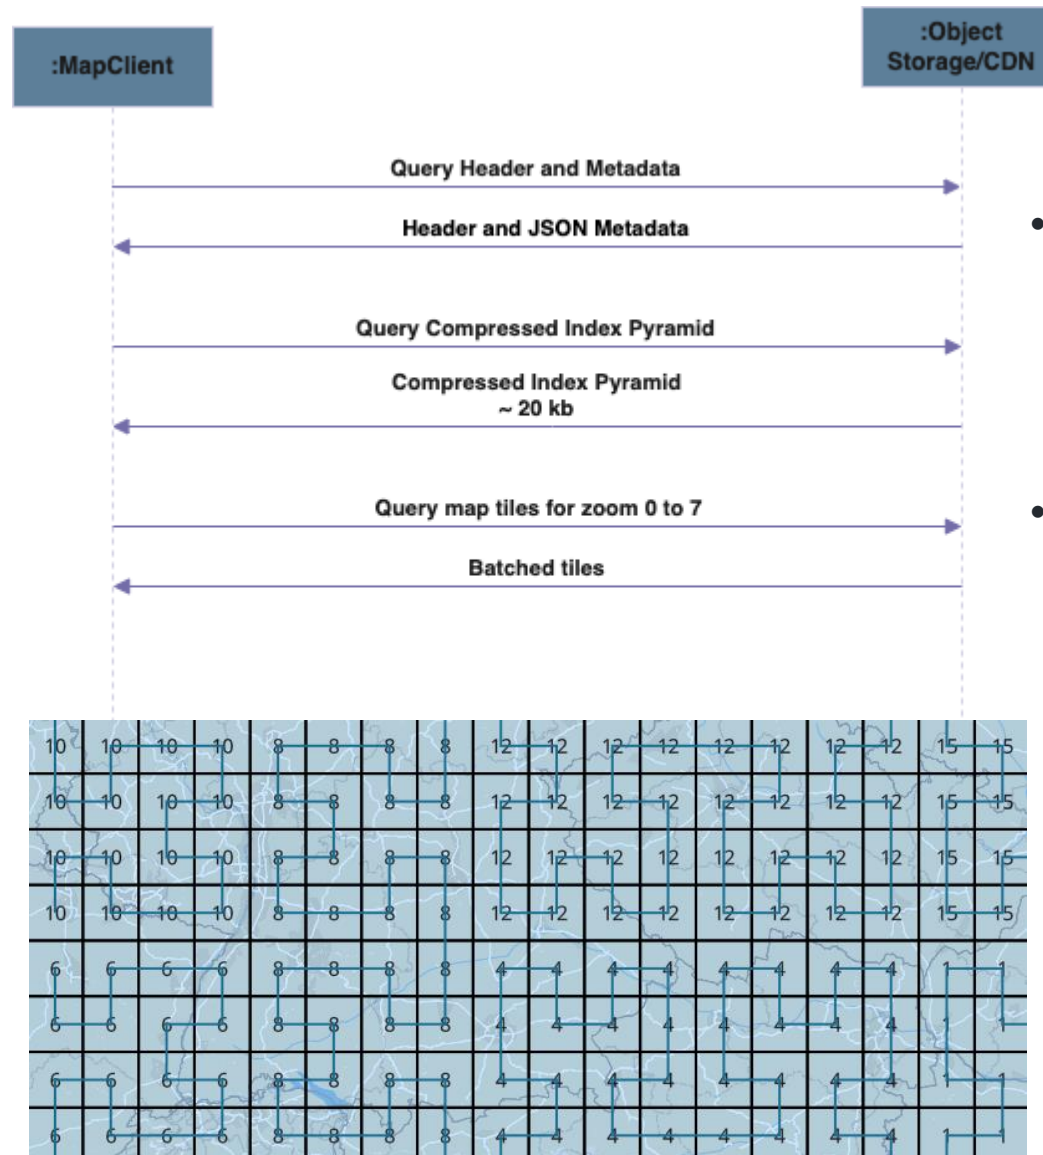
Load overview index records for exploring the map

→ ~ 20k index records for zoom 0 to 7 for a planet tileset

# COMTiles sequence

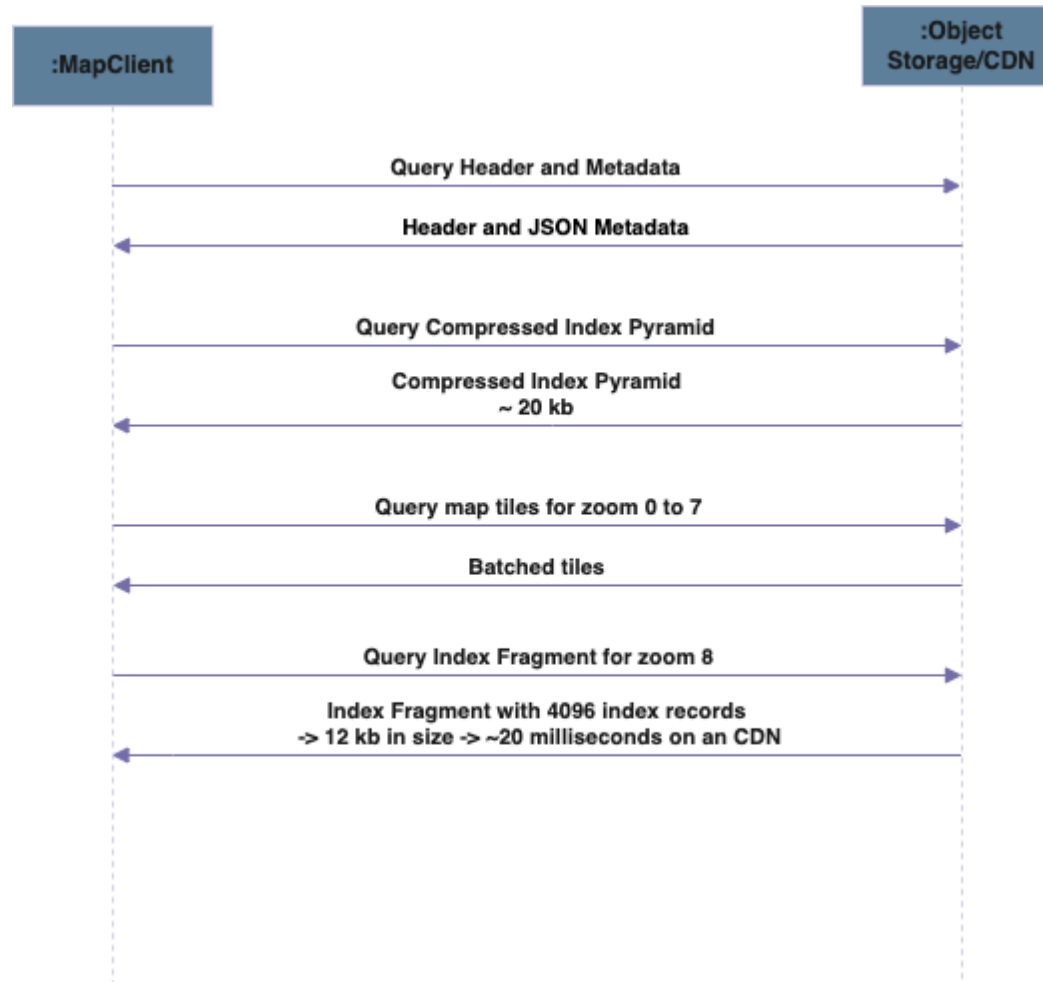


# COMTiles sequence



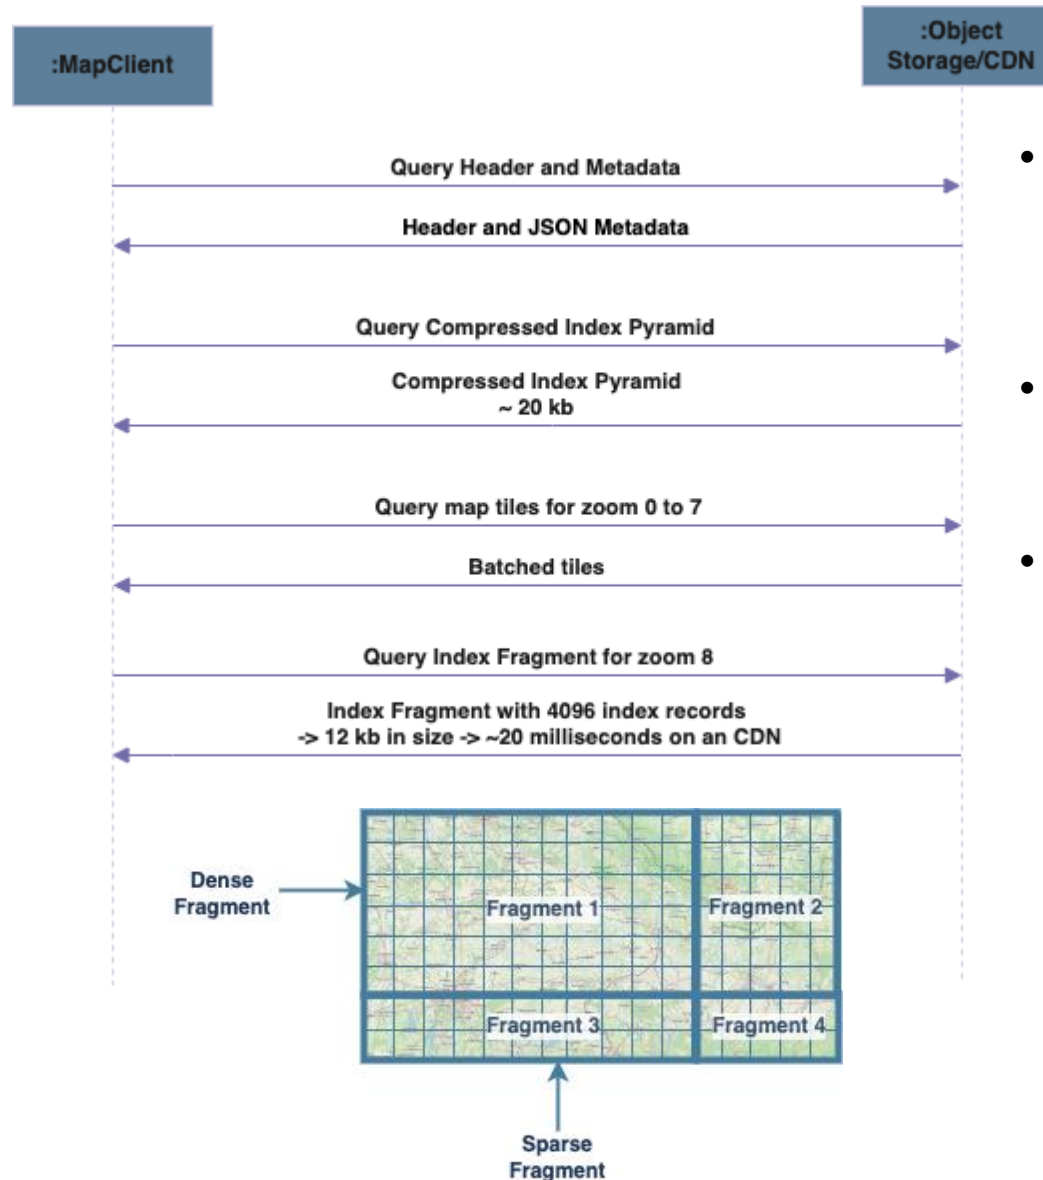
- The individual tile requests can be batched to improve performance (for HTTP/1.1 requests) and in particular to reduce the transfer costs
- This can reduce the number of tile requests by up to 90%

# COMTiles sequence



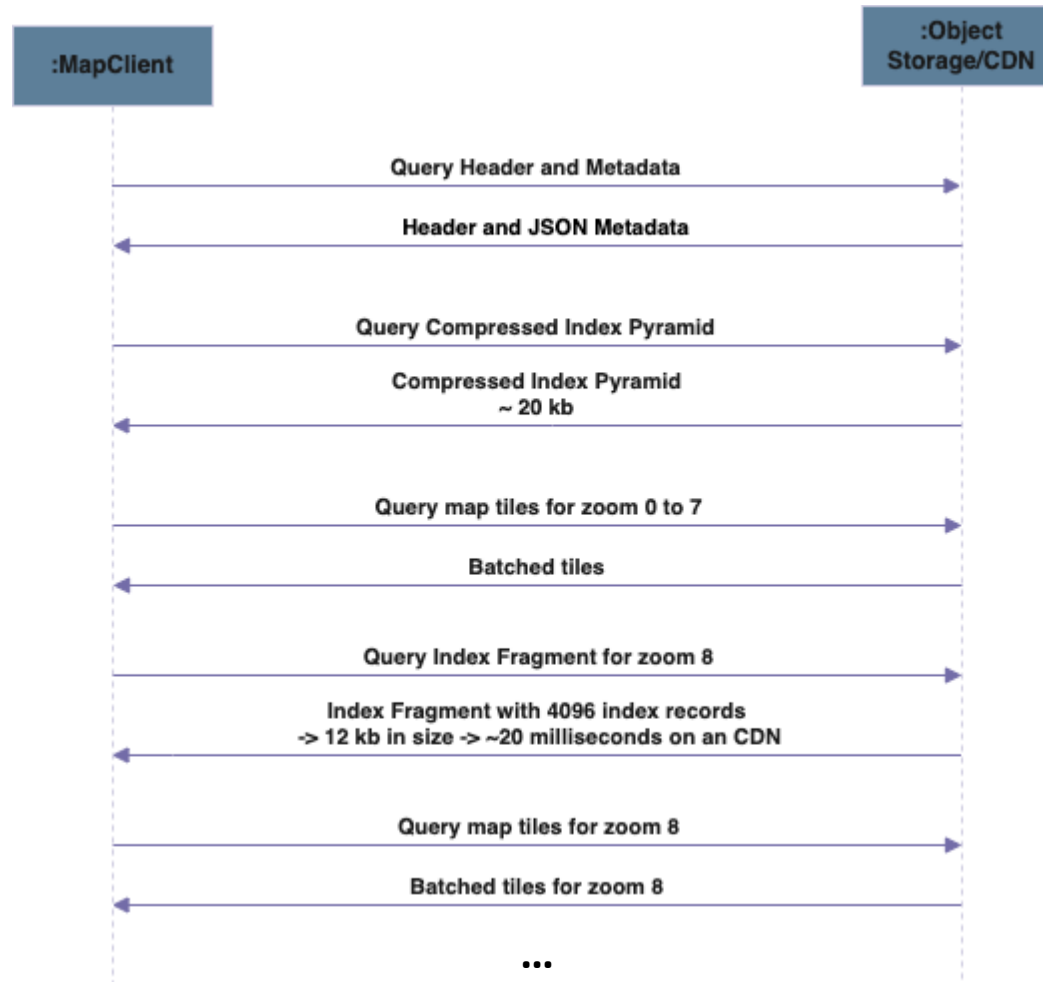


# COMTiles sequence



- For higher zoom levels portions of the index are lazy loaded via **index fragments**
- The index fragments are ordered on a **space-filling curve** (Hilbert or Z-Order)
- Most of the time only one additional pre-fetch per zoom level is needed before accessing the actual map tiles for the current viewport of the map -> **this is bearly or not at all noticeable for the user regarding the user experience**

# COMTiles sequence



# Challenges

- Additional requests for fetching portions of the index → a “servless” Tile Server (e.g. hosted on AWS Lambda or CloudFlare Workers) can be used to eliminate the additional index request and to further minimize the latency
- Storing large tilesets on an object storage is cheap, but egress costs can get very expensive → with “requester pays” the users have to pay for the egress costs
- Multipart ranges are not supported by the cloud providers on object storage only on a CDN
- Most cloud provider doesn’t support HTTP/2 on object storage only on a CDN
- Compression via Content-Encoding header is not supported in http range requests → Compression Streams API also not (yet) supported on all browsers
- Automatic (dynamic) compression is supported for regular HTTP requests by the CDNs but not for range requests
- Readoptimized format → updates are expensive

# Next Steps

- Cloud optimized 3D Tiles
  - The 3DTiles Next implicit tiling extension has now support for fixed spatial data structures (quadtree for 2D or octree for 3D) which allow random access
  - Combining the COMTiles approach with 3D Tiles Next for building a cloud optimized 3D geospatial format
- Cloud optimized vector tiles
  - Column-oriented approach like Parquet for the properties of the features to achieve better compression
  - Faster decoding (e.g. zero copy) inspired by GLTF (3DTiles)
  - Make a specific layer of a tile queryable via additional secondary index (sidecar) → can reduce the number of features up to 50% depending on the style

# Thank You!

Email:  
markus.tremmel23@gmail.com

GitHub: @mactrem